FACULTY
OF MATHEMATICS
AND PHYSICS
Charles University

MASTER THESIS

Vojtěch Hudeček

Exploiting user's feedback to improve pronunciation of TTS systems

Institute of Formal and Applied Linguistics
Faculty of Mathematics and Physics
Charles University in Prague

Supervisor of the master thesis:	doc. Ing. Zdeněk Žabokrtský, Ph.D.
Study programme:	Informatics
Study branch:	Artificial Intelligence

Prague 2017

I declare that I carried out this master thesis independently, and only with the cited sources, literature and other professional sources.

I understand that my work relates to the rights and obligations under the Act No. 121/2000 Sb., the Copyright Act, as amended, in particular the fact that the Charles University has the right to conclude a license agreement on the use of this work as a school work pursuant to Section 60 subsection 1 of the Copyright Act.

In .... date ....   signature of the author

Title: Exploiting user's feedback to improve pronunciation of TTS systems

Author: Bc. Vojtěch Hudeček

Department: Name of the department

Supervisor: doc. Ing. Zdeněk Žabokrtský, Ph.D, Institute of Formal and Applied Linguistics

Abstract: Although spoken dialogue systems have greatly improved, they still cannot handle communications involving unknown topics and are very fragile. We will investigate methods that can improve spoken dialogue systems by correcting or even learn the pronunciation of unknown words. Thus we will provide better user experience, since for example mispronounced proper nouns are highly undesirable. Incorrect pronunciation is caused by imperfect phonetic representation, typically phonetic dictionary. We aim to detect incorrectly pronounced words by exploiting userâĂŹs feedback as well as using prior knowledge of the pronunciation and correct the transcriptions accordingly. Furthermore, the learned phonetic transcriptions can be used to improve speech recognition module by refining its models. Models used in speech recognition cannot handle words that are not in their vocabulary or have phonetic representation. Extracting those words from userâĂŹs utterances and adding them to the vocabulary should lead to a better overall performance.

Dedication.

Contents

1. Introduction

Voice communication is a common interface of many systems nowadays. Voice communication applications range from simple one-word control commands to complex communication in spoken dialogue systems. In this work, we consider mainly such complex systems. We now briefly introduce setting of such system.

Leading a reasonable dialogue with a user is a substantially difficult composed of many subtasks. High level overview of such system is given in Figure 1. First, the input from the user has to be processed Spoken dialogue systems typically contain Automatic Speech Recognition (ASR) module for this purposes, so the natural speech can be recognized and translated into words. Afterwards, some Spoken Language Understanding (SLU) technique has to be employed in order to extract essential information from the input. In plain words, we have to understand the meaning of the utterance obtained from the user. The dialogue system than derives a relevant response using its policy. The process of this derivation depends on particular implementation. A $rule-based$ policy can be used, however it is more common nowadays to employ $data-driven$ approaches.

The derived response has to be translated into human readable language, using Natural Language Generation (NLG) techniques. The response can be displayed in textual form, however, it is more common to generate audio recording with human voice reading the response. Although it is possible to use a set of prerecorded utterances, this approach has obvious limitations since it is not able to read an arbitrary phrase. Particularly, it may be difficult to read named entities and numerical values such as time and date. Also, the usage of variable utterances provides better user experience.
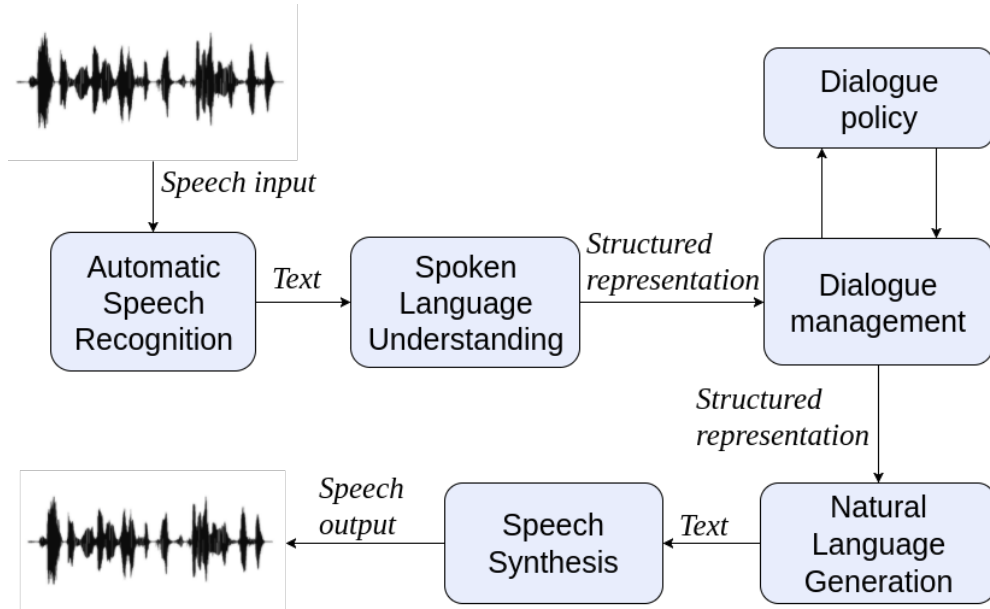

Figure 1.1: A high level architecture of spoken dialogue system.

To introduce a variability and allow to use nearly arbitrary phrase a Text-To-Speech (TTS) module is usually also part of dialogue systems. The purpose of this module is to transform written text utterance to natural speech. Modern TTS

systems produce audio waveforms that sound naturaly and the pronunciation is sufficiently good. Nevertheless, it may experience some difficulties, mainly when it comes to unknown words. This may happen, because although some multilingual TTS systems exists ([**?**], [**?**]), the system is usually trained using certain set of words, typically from one language. But some applications often require to pronounce named entities or other language- or domain-specific words, that cannot be present during the training phase. Furthermore, the system needs to obtain information, which language it should use. This does not make much sense when considering individual proper names. Such words are called Out-Of-Vocabulary (OOV), because the TTS systems are vocabulary based, i.e. they contain set of known words and their phonetic transcriptions. Occurences of OOV words cause situations, when the system has to employ some mechanism to derive the pronunciation and the words may be mispronounced. The derivation of the pronunciation is inherently imperfect and may cause errors. Although this does not occur often, the negative effect can be quite strong, since it is inconvenient for the user, especially when known propername, e.g. his or her name is pronounced with mistakes.

In this work, we aim to improve the TTS system pronunciation of OOV words. First, we explore methods that can identify words that are potentially difficult to pronounce. The identification is first step towards the improvement, because we can then ask the user for proper pronunciation. We propose and explore several measures that can reflect badly pronounced words without any prior language knowledge.

Next, we try to improve the TTS system's pronuncication of the desired words. To achieve this, we ask the user and obtain correct pronunciations from him. So we get training examples and we are able to improve the TTS system by processing the obtained recording, deriving a phonetic transcription (i.e. pronunciation) and adding it to the TTS vocabulary. Moreover, the derived pronunciations can be used to improve the recognition ability of the ASR module, since it is also dictionary-based. Thus if we enlarge the vocabulary, the ASR has got access to more possibilities and can recognize the new words correctly.

There are several issues, that are related with the OOV words so methods of deriving correct pronunciations has got potentially very useful applications. As it has been said, it can be used to enlarge vocabularies of TTS or ASR systems. This may potentially lead to better pronunciation in case of TTS. It is because the extended vocabulary implies, that a number of occurences of the badly pronounced words decreases. In case of ASR systems, bigger vocabulary can cause, that less words are recognized incorrectly. There exist several ways how to obtain such feedback, however, this is not a subject of this work. It is related with respective dialogue policy and concrete applications. Theoretically, the method can work with just one gold example, however, it is better to obtain more recordings in general. In Figure 1.2 we provide basic example of simple dialogue, illustrating how real application could look like. However, in this work we assume the user's recording(s) have been gathered already and we do not discuss the dialogue policy.

```
System:  Hello, /AANDRZHEZH/.
User:  You said it wrong, my name is /ONDRZHEI/.
System:  /ANDREY/, correct?
User:  No, it is /ONDRZHEI/.
System:  Oh, /ONDRZHEI/?
User:  That's right.
```

Figure 1.2: Sample dialogue illustrating the pronunciation correction. The transcriptions of the user's name are given in ARPABET[1]

2. Overview of used techniques and algorithms

2.1 Audio signal processing

We sketch here the basic principles as described in [13], so we can understand our input data. Speech signal in the real world are mechanical waves, so it is obviously analogous quantity. When we process the speech signal with computers, it is assumed to be digitised, so it is converted into discrete form. When performing digital signal processing, we are usually concerned with three key issues.

1. To remove the influence of phase, because the ear is not sensitive to phase information in speech, and we can use frequency domain representation.

2. Performing source/filter separation, so that we can study the *spectral envelope* of sounds. Spectral envelope characterizes the frequency spectrum of the signal, which is essential for the speech recogniton and processing.

3. We often wish to transform these spectral envelopes and source signals into more efficient representations.

Overview of the main standard processing steps follows.

The input signal is divided into *windows*. Windowing considers only some part of the signal. Usually, the signal is transformed at window borders to prevent discontinuities. This is achieved for example by use of *Hamming* window. The complete waveform is therefore a series of windows, that are sometimes called *frames*. The frames overlap, this is influenced by the size of the *frame shift*. Inside one frame, we assume, that the speech signal is stationary and we transform it to a frequency domain by Discrete Fourier Transformation[1].

It is better for observing the natural speech characteristics, to use a logarithm scale instead of linear. In fact, the *mel scale*[2] is frequently used, which even better corresponds with the human perception. Further, the so called *cepstrum*[3] is computed from the log magnitude spectrum, by performing inverse Fourier transformation. Cepstrum can be represented by coefficients, number of which can be chosen. Those are called the *Mel Frequency Cepstral Coefficients(MFCCs)*. To sum up, we have described how to process the digital audio signal and convert it into discrete series of overlaping frames, each of which is represented by a fixed-length vector of MFCCs.

We gave a brief overview of the signal processing procedure that is essential to work with the audio signal. When we mention the input speech signal, we mean series of MFCC vectors, unless explicitly stated otherwise.

---

[1]https://en.wikipedia.org/wiki/Discrete_Fourier_transform
[2]https://en.wikipedia.org/wiki/Mel_scale
[3]https://en.wikipedia.org/wiki/Cepstrum

2.2 Mel Cepstral distortion

Mel Cepstral Distortion [**?**] is a well described measure, that should mirror differences between speech samples. There are some caveats worth stating at the outset. Firs, there are many other factors that contribute to the perception of voice quality. For example it takes no account of speech dynamics, either short-range differentials or long-range prosodic effects. Morevoer, distortions in the pitch contour are ignored. However, it should reflect the differences in speech quality and pronunciation.

As we have described in the previous section, the audio sample can be described by its *cepstrum*. This cepstrum may be then represented by Mel Frequency Cepstral Coefficients (MFCCs). Its order can be chosen arbitrarily, we used 35 coefficients length vectors. Thus we can transform each audio record into a sequence of float vectors with length 35. The MCD [**?**] of two sequences is basically normalized Mean Squared Error between those sequences after aligning. We define it as follows:

$$MCD(v^1, v^2) = \frac{\alpha}{T'} \sum_{ph(t) \neq SIL} \sqrt{\sum_{d=1}^{D}(v_d^1(t), v_d^2(t))^2} \qquad (2.1)$$

The T' stands for the number of frames of the shorter of the recordings. The common use of MCD involves omitting the 0-th coefficient, since it corresponds to the intensity of signal and therefore it is not usually desired. The scaling factor α is present   for   historical   reasons.

The MCD measure has one substantial disadvantage: it does not reflect, that the two compared recordings may not have the same length and, what is more, they can be misaligned. Thus, we use a modified algorithm that first aligns the sequences using Dynamic Time Warping and then computes MCD with the result.

2.3 Finite state transducers

Finite state transducers (FSTs) [7], are basically finite state automata that are augmented by addition of output labels to each transition. We can further modify the FST and add a weight to each edge. We define a semiring $(S, \oplus, \otimes, \overline{0}, \overline{1})$ as a system that fullfills the following conditions:

- $(S, \oplus, \overline{0})$ is a commutative monoid with identity element $\overline{0}$

- $(S, \otimes, \overline{1})$ is a monoid with identity element $\overline{0}$

- $\otimes$ distributes over $\oplus$

- $\forall a \in S : a \otimes \overline{0} = \overline{0} \otimes a = \overline{0}$

With this definition of the semiring, we can describe a weighted FST [8] T over a semiring S formally as a 8-tuple $(\Sigma, \Delta, Q, I, F, E, \lambda, \rho)$ provided that:

- Σ, Δ are finite input and output alphabets

- Q is a finite set of states

- $I \subset Q$ is a set of initial states

- $F \subset Q$ is a set of final states

- E is a multiset of transitions

- $\lambda : I \to S$ is an initial weight function

- $\rho : F \to S$ is a final weight function.

Each transition is element of $Q \times \Sigma \times \Delta \times S \times Q$. Note, that since E is a multiset, it allows two transitions between states p and q with the same input and output label, and even the same weight. However, this is not used in practice. An example of weighted FST is given in Figure 2.3.

Finite state transducers have been used widely in many areas, especially linguistics. They can represent local phenomena encountered in the study of language and they usage often leads to compact represantations. Moreover, it is also very good from the computational point of view, since it advantageous in terms of time and space efficiency. Whole area of algorithms has been described that consider FSTs. One of the most important is *determinization*, which determinize paths in the FST and thus allows the time complexity to be linear in the length of input. Weighted FSTs have been widely used in the area of automatic speech recogniton. The allows to compactly represent the hypothesis state of the acoustic model as well as combining it with the information contained in the language model.
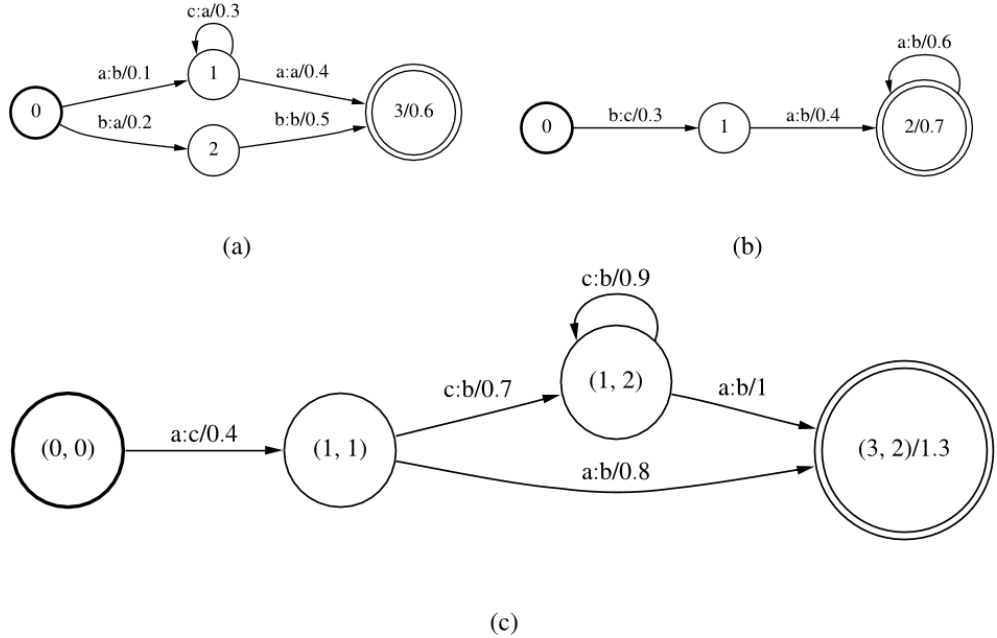

Figure 2.1: Example of Finite State Transducers and their composition

2.4 Automatic speech recognition (ASR)

Overview

The task in ASR is quite clear: An utterance spoken in natural language should be translated into its text representation. Formally, we are given a sequence of acoustic observations $\mathbf{X} = X_1 X_2 \ldots X_n$ and we want to find out the corresponding word sequence $\mathbf{W} = W_1 W_2 \ldots W_m$ that has a maximum posterior probability $P(W|X)$i.e., according to the Bayes rule:

$$\hat{\mathbf{W}} = argmax_w \frac{P(\mathbf{W})P(\mathbf{W}|\mathbf{X})}{P(\mathbf{X})} \qquad (2.2)$$

The problem's solution consists of two subtasks. First, we have to build accurate acoustic model that describes the conditional distribution $P(W|X)$. The second problem is to create a language model that reflects the spoken language that should be recognized. Usually, a variation of the standard N-gram approach is sufficient.

Individual words are composed of phonetic units, which are in turn modeled using states in probabilistic machinery such as a finite state transducers. Because the speech signal is continuous, it is transformed to discrete sequence of samples. MFFC's are commonly used to represent the signal. The process of deriving this sequence is described in further detail in section 2.1. In general, it is difficult to use whole-word models for the acoustic part, because every new task may contain unseen words. Even if we had sufficient number of examples to cover all the words, the data would be too large. Thus, we have to select basic units to represent salient acoustic and phonetic information. These units should be *accurate*, *trainable* and *generalizable*. It turns out that the most appropriate units are phones. Phones are very good for training and generalization, however, they do not include contextual information. Because of this, so called triphones are   commonly   used.

To model the acoustics, Hidden Markov Models(HMM) are commonly used, because they can deal with unknown alignments. In recent years, neural networks have experienced big breakthrough and they found application even in the area of speech recognition. Namely, recurrent neural networks are able to work with an abstraction of memory and process sequences of variable length, so it overcomes the HMM's in terms of accuracy. The language and acoustic models are traditionally trained independently and they are joined during the decoding process. The decoding process of finding the best matched word sequence $\mathbf{W}$ to match the input speech signal $\mathbf{X}$ in speech recognition systems is more than a simple pattern recognition problem, since there is an infinite number of word patterns to search in continuous speech recognition. A lattice is built in order to make the decoding. Its nodes represent acoustic units and edges are assigned costs. This assignment corresponds to the likelihood determined by the acoustic model as well as the language model, so the FST is a composition of theses two. Decoding the output is realized as searching the best path through this lattice.

The decoding graph is costructed in such a way that the path can contain only

words present in the vocabulary the language model is built on. Thus, each possible path consists of valid words. Because of that, we can see the word as a basic unsplittable unit, although it is actually composed from phones, or triphones respectively. However, we can build an FST that allows to construct the path from phones and thus recognize the input on the phonetic level. Also the search algorithm preserves alternative paths so in the end it gives us not only the best path but also list of alternative hypotheses. We call it the n-best list. Each hypothesis in the n-best list is associated with its likelihood that expresses information from both the language and acoustic model. In TODO:chapter we explore the n-best list and propose method that exploits it.

2.5   Text-to-speech (TTS)

Overview

[13] The text-to-speech problem can be looked at as a task of transforming an arbitrary utterance in natural language from its written form to the spoken one. In general, these two forms have commonalities in the sense that if we can decode the form from the written signal, then we have virtually all the information we require to generate a spoken signal. Importantly, it is not necessary to (fully) uncover the meaning of the written signal, i.e. employ Spoken Language Understanding (SLU).

On the other hand, we may need to generate prosody which adds some information about emotional state of the speaker or emphasizes certain parts of the sentence and thus changing the meaning slightly. To obtain prosodic information, sophisticated techniques need to be involved, since its not a part of common written text, except some punctuation. Another difficulties stem from the fact, that we often need to read numbers, or characters with special meanings such as dates or mathematical equations. The problem of converting text into speech has been heavily explored in the past and the TTS systems (engines) are very sophisticated nowadays. The subsequent section, describes briefly the typical architecture of TTS systems.

TTS system architecture

Let us now describe a common use of TTS system. The architecture usually consists of several modules, although some end-to-end systems base on neural networks have appeared recently ([15], [16]). First, the input text is divided into sentences and each sentence is further tokenized, based on whitespace characters, punctuation etc. Then, the non-natural language tokens are decoded and transformed into text form. We then try to get rid of ambiguity and do prosodic analysis. Although much of the information is missing from the text, we use algorithms to determine the phrasing, prominence patterns and tuning the intonation of the utterance.

Next is the synthesis phase. The first stage in the synthesis phase is to take the words we have just found and encode them as phonemes. There are two

main approaches how to deal with the synthesis phase. More traditional approach is so called **unit concatenation**. This approach uses a database of short prerecorded segments of speech (roughly 3 segments per phoneme). These segments are then concatenated together with some use of signal processing so they fit together well. An alternative is to use machine learning techniques to infer the specification-to-parameter mapping from data. While this and the concatenative approach can both be described as data-driven, in the concatenative approach we are effectively memorizing the data, whereas in the statistical approach we are attempting to learn general properties of the data.

Trained models are able to transform the phonemes into audio signal representation. This representation is then synthesized into waveforms using a vocoder module. Although unit concatenation approaches generally achieve better results, the statistical ones are more flexible and have good possibilities to fine tuning or postprocessing.

Grapheme to Phoneme conversion and its drawbacks

In the stage of converting graphemes (i.e. text) to phonemes, a *g2p* module is commonly used. The g2p task is to derive pronunciations from ortographic transcription. Traditionally, it was solved using decision trees models. It can also be formulated as a problem of translating one sequence to another, so neural networks can be used to solve this task ([17]). However, in our work we use joint-sequence model ([2]), which estimates joint distribution of phonemes based on probabilities derived from *n*-grams. The *g2p* module is a crucial component of the system and it has great impact on the final pronunciation. Since we use machine learning techniques and train on the dictionary data, the model inherently adopts pronunciation rules from the respective language. Although it is in principle possible to train multilingual *g2p* ([11]), it is usually not the case in common TTS systems. The problem arises when words that originate in some foreign language should be pronounced. In this work we address this issue by deriving pronunciations from the ASR output.

Speech Synthesis Markup Language (SSML)

[14] SSML is an XML standard that allows to specify aspects of the speech. It is an input to the synthesis module and many TTS engines support its use. We can specify emotions, breaks etc. with the help of SSML. More importantly, it can be used to input particular phonemes and thus circumvent the *g2p* module. In our work we use this method to feed our phonetic transcriptions into the engine. The disadvantage of this method is, that many TTS engines process the SSML input imperfectly or not at all. However, in principle it is possible to add the transcriptions directly into the system's vocabulary.

```
<?xml version="1.0"?> <speak version="1.0" xmlns="..."
xmlns:xsi="..." xsi:schemaLocation="..." xml:lang="en-US">
    <voice gender="female">Mary had a little lamb,</voice>
    <voice gender="female" variant="2">
        Its fleece was white as snow.
    </voice>
    <voice name="Mike">I want to be like Mike.</voice>
</speak>
```

Figure 2.2: Example of SSML code

Overview of the used TTS systems

1. Cereproc[4] This engine is a commercial software that is free for educational purposes. Although it is not open-sourced, we use it for its good quality and reliable outputs.

2. MaryTTS[5] MaryTTS is a German open-source project written in Java. It is highly customizable and modular. Thus we can explore output of an arbitrary module or replace it in the processing pipeline. MaryTTS is based on client-server architecture.

3. gTTS[6] is a TTS service provided online by Google. It achieves very good quality, however it allows ony use of one voice in the free version.

4. Pico [7] SVOX Pico TTS is a lightweight engine that lacks a good quality of the gTTS, however it offers a large selection of voices and it is commonly used on the phones with Google's Android operating system.

2.6 Algorithms description

2.6.1 Dynamic Time Warping (DTW)

[9] The measurement of similarity between two time series is an important component of many applications. Moreover, sometimes we need to align two sequences that describe same data but are of different lengths. Both this tasks can be solved with use of DTW. Suppose we have two time series, a sequence Q of length n, and a sequence C of length m, where

$$Q = q_1, q_2, \ldots, q_i, \ldots, q_n$$
$$C = c_1, c_2, \ldots, c_j \ldots, c_m$$

To align these two sequences using DTW, we first construct an n-by-m matrix where the (i^{th}, j^{th}) element of the matrix corresponds to the squared distance, $d(q_i, c_j) = (q_i - c_j)^2$, which is the alignment between points q_i and c_j. To find the best match between these two sequences, we retrieve a path through the matrix

---

[4]https://www.cereproc.com/
[5]http://mary.dfki.de/
[6]https://pypi.python.org/pypi/gTTS
[7]https://github.com/stevenmirabito/asterisk-picotts

that minimizes the total cumulative distance between them. In particular, the optimal path is the path that minimizes the warping cost:

$$DTW(Q, C) = \sqrt{\sum_{k=1}^{K} w_k} \qquad (2.3)$$

where w_k is the matrix element $(i, j)_k$ that also belongs to k^{th} element of a warping path W, a contiguous set of matrix elements that represent a mapping between Q and C. Methods of dynamic programming are used to fill in the values and find alignment of sequences. Thus a result of this algorithm is a mapping σ that can be used to construct sequence of pairs A containing members of both sequences C, Q in each time step. σ is constructed in such a way, that it holds:

$$A_i = (C_{\sigma(C,i)}, Q_{\sigma(Q,i)}); i = 1 \ldots K \qquad (2.4)$$

$$max(|C|, |Q|) \leq K \qquad (2.5)$$

$$\sum_{i=0}^{K} d(\sigma(C, i), \sigma(Q, i)) \, is \, minimal \qquad (2.6)$$

Where $d(x, y)$ is an arbitrary distance metric. In our application, we want to align two sequences of phonemes, hence we can treat it as strings and work with repsective distance. We choose Levensthein distance to be our metric.

2.6.2 Clustering

Cluster analysis or clustering [8] is the task of grouping a set of objects in such a way that objects in the same group (called a cluster) are more similar to each other than to those in other clusters. The measure of similarity can be arbitrary distance measure. Usually, euclidean distance is used when considering points in space, however it is not a requirement. Clustering is inherently an unsupervised method, we often want to describe unknown data, that is find some structure or local similarities.

There exist different methods of clustering. *Hierarchical clustering* is based on the core idea of objects being more related to nearby objects than to objects farther away. These algorithms connect object" to form clusters based on their distance. Hierarchical clustering is a bottom-up method, meaning that it forms clusters iteratively. The hierarchical clustering algorithms' result can be represented as a dendrogram, which is a tree diagram representing relations among the explored data. An example of dendrogram is given in Figure 2.6.2. We can make a horizontal cut to derive respective clusters.

On the other hand in *Centroid-based clustering*, clusters are represented by a central vector, which may not necessarily be a member of the data set. If we fix the number of clusters to k, we can see the task as an optimization problem: find the k cluster centers and assign the objects to the nearest cluster center, such that the squared distances from the cluster are minimized. Typical algorithm

---

[8]https://en.wikipedia.org/wiki/Cluster_analysis

Figure 2.3: An example of dendrogram outputted by hierarchical clustering methods.

from this group is k-means clustering[9]. which is an Expectation-Maximization (EM) algorithm.

Spectral clustering [**?**] The task of finding good clusters has been the focus of considerable research in machine learning and pattern recognition. Often, EM algorithms are used to learn a mixture density. Unfortunately, these aproaches have several drawbacks. First, to use parametric density estimators, simplifying assumptions usually need to be made (e.g., that the density of each cluster is Gaussian). Second, the log likelihood can have many local minima and therefore multiple restarts are required to find a good solution using iterative algorithms. An alternative is to use spectral methods for clustering. Here, one uses the top eigenvectors of a matrix derived from the distance between points. Therefore, the mentioned problems are overcomed.

---

[9]https://en.wikipedia.org/wiki/K-means_clustering

2.6.3   Basic notions in speech communication

Communication is a process during which some entities exchange information. There are different types of communication, we focus here on the communication in natural language. When we consider communication, we should define what we understand by a term form. This concept deserves some elaboration. A proper understanding of this part of communication is essential for our purposes, but thankfully we can rely on established frameworks.

The main issue when discussing form, or a message, is to distinguish the general nature of the form from any specific instantiation it may have. Imagine, someone wants to deliver an information by showing some color. If a red card is used, it isn't the card itself that is the important point, rather than the red color was used instead of different color. So for example, a different red card could be used successfully as it is not the particular red card itself which is the form, but rather that the card has the property "red" and it is this property and its contrasting values that is the basis of the communication system. We use the term signal to refer to the physical manifestation of the form. Here the term signal is again used in a specific technical sense: it can be thought of as the thing which is stored or the thing which is transmitted during communication. The idea of "red" and the fact that it contrasts with other colors is the form; the physical light waves that travel are the signal.

It is important to see that the relationship between meaning and form, and form and signal are quite separate. If we assigned meaning only to colors red, blue and green and an orange card would have been showed to us, we would not decode the meaning, although the signal-form translation worked well. This situation is similar to that when we read a new word - while are quite capable of using our eyes to read the physical pattern from the page and create an idea of the letters in our head, if we don't know the form-meaning mapping of the word, we can not understand it. Finally, we use the term channel or medium to refer to the means by which the message is converted and transmitted as a signal.

In this work, we deal mainly with the spoken channel and the written channel. Encoding is the process of creating a signal from a message. When dealing with speech, we talk of speech encoding and when dealing with writing we talk of writing encoding. Speech encoding by computer is more commonly known as speech synthesis, which is of course the topic of this book. The most significant aspect of speech encoding is that the nature of the two representations, the message and the speech signal, are dramatically different. A word (e.g. $HELLO$) may be composed of four phonemes ($/H\ EH\ L\ OW/$) but the speech signal is a continuously varying acoustic waveform, with no discrete components. In this work, we focus on derivation of the first representation and we let the used synthesis frameworks do the speech signal generation. So we simplify the speech representation to its phonetic transcription.

However, it is important to understand, that we do not change the meaning or even form of the communication, we just want to use different channel for the transmition. When we consider the verbal language, we may talk about the duality of form such that we have phonemes which combine to words which combine to form sentences. We define three terms that help us to describe the problematics:

1. **phonemes** are members of the relatively small set of units which can be combined to produce distinct word forms.

2. Term **phone** is used to describe a single speech sound, spoken with a single articulatory configuration.

3. **allophones** link the phonemes and phones: different ways of realising a single phoneme are called allophones.

The duality principle states that verbal language is not one system but two; the first system uses a small inventory of forms, phonemes, which sequentially combine to form a much larger inventory of words.

We have a similar dual structure in written communication, where words are made of **graphemes**. Graphemes are in many ways analogous to phonemes, but differ in that they vary much more from language to language. In alphabetic languages like English, graphemes can be thought of as letters or characters; in syllabic writing like Japanese hiragana, each grapheme represents a syllable or mora (unit that determines syllable weight, i.e. stress or accent) and in languages like Chinese each grapheme represents a full word or part of a word. We consider alphabetic languages in this work, encoded with the Latin alphabet[10]. However, it is possible to extend the used grapheme-to-phoneme conversion technique [2] to work with other types of languages.

To encode the phonemes, phonetic alphabets are used. The standard is an International Phonetic Alphabet (IPA)[11]. IPA symbols are composed of one or more elements of two basic types, letters and diacritics. The general principle of the IPA is to provide one letter for each distinctive sound (speech segment), although this practice is not followed if the sound itself is complex. An example of IPA characters is given in Figure 2.6.3.

The problem is that some of the tools we work with use different phonetic alphabets with possibly different rules. However, the TTS engines work with IPA, so it is neccessary to convert the transcriptions. To overcome this problem, we have derived one-to-one mappings between these sets and IPA. This process may introduce some errors, however the alphabets are usually based on similar principles, so it is no so hard to derive this mappings. Namely, we used the Arpabet[12], Timit[13] and CGN[14] phoneme sets.

As we have suggested, some speech synthezizers are able to accept input in the SSML format which we introduce in 2.5. However, not all the synthesizers have this ability. We use Cereproc to work with utterances written in SSML. An example of such utterance is given in Figure 2.6.3. The problem with providing transcriptions in this way is, that the Cereproc synthesizer ignores the characters defining strech of the phoneme. Therefore, we ommit these characters from the transcriptions by mapping the streched phonemes to their base variants This leads to worse quality of the pronunciation, however is sufficient to explore if the transcription contains sensible phonemes.

---

[10]https://en.wikipedia.org/wiki/Latin_alphabet

[11]https://en.wikipedia.org/wiki/International_Phonetic_Alphabet

[12]https://en.wikipedia.org/wiki/Arpabet

[13]https://catalog.ldc.upenn.edu/docs/LDC93S1/PHONCODE.TXT

[14]https://pdfs.semanticscholar.org/64d5/d612f9c639271a1340820a418fa5baf02770.pdf

Figure 2.4: Some letter of the IPA alphabet together with their occurences in English words.

...

```
<phoneme alphabet="ipa" ph="təmeiɾou"> tomato </phoneme>
```

...

Figure 2.5: An example of the synthesizer's input in SSML, forcing it to use our phonetic transcription.

2.7 Related Work

2.7.1 Grapheme-to-phoneme conversion

An automatic grapheme-to-phoneme conversion was first considered in the context of TTS applications. The input text needs to be converted to a sequence of phonemes which is then fed into a speech synthesizer. It is common in TTS systems that they first try to find the desired word in the dictionary and if it doesn't find it, it employs the grapheme-to-phoneme ($g2p$) module. A trivial approach is to employ a *dictionary look-up*. However, it cannot handle context and inherently covers only finite set of combinations. To overcome this limitations, the rule-based conversion was developed. Kaplan and Kay [4] formulate these rules in terms of finite-state automata. This system allows to greatly improve coverage. However the process of designing sufficient set of rules is difficult, mainly since it must capture irregularities. Because of this, a *data-driven* approach based on machine learning has to be employed. Many such techniques were explored, starting with Sejnowski and Rosenberg [12]. The approaches can be divided into three groups.

Techniques based on local similarities

The techniques presuppose an alignment in the training data between letters and phonemes or create such an alignment in a separate preprocessing step. The alignment is typically construed so that each alignment item comprises exactly one letter. Each slot is then classified (using its context) and a correct phoneme is predicted. Neural networks and decision tree classifiers are commonly used for this task.

Pronunciation by analogy

This term is typically used for methods that could be described as nearest-neighbor-like. They search for local similarities in the training lexicon and the output pronunciation is chosen to be analogous to retrieved examples. In the work of Dedina and Nusbaum [3] the authors first identify words from the database that match the input string and then build a pronunciation lattice from them. Paths through the lattice then represent the derived pronunciations.

Probabilistic approaches

The problem can also be viewed from a probabilistic perspective. Pioneers in this area were Lucassen and Mercer [5]. They create $1 - to - n$ alignments of the training data using a context independent channel model. The prediction of the next phoneme is based on a symmetric window of letters and left-sided window of phonemes. Authors then construct regression tree, the leafs of which carry probability distribution over phonemes.

A popular approach based on probability modelling is to employ so called *joint sequence models* [2]. We use an open-source implementation of such a model in our work. This approach formalizes the task as follows:

$$\boldsymbol{\varphi}(\boldsymbol{g}) = \operatorname*{argmax}_{\varphi' \in \Phi^*} p(\boldsymbol{g}, \boldsymbol{\varphi}') \tag{2.7}$$

where * denotes a Kleene star. In other words, for a given ortographic form $\boldsymbol{g} \in G^*$ we want to find the most likely pronunciation $\boldsymbol{\varphi} \in \Phi^*$, where G, Φ are ortographic and phonetic alphabets. The fundamental idea of joint-sequence models is that the relation of input and output sequences can be generated from a common sequence of joint units. These units carry both input and output symbols. Formally, the unit called *grapheme* is a pair $q = (\mathbf{g}, \varphi) \in Q \subset G * \times \Phi$ where $\mathbf{g}$ and φ are letter and phoneme sequences which can be of different lengths. In the simplest case, each unit carries zero or one input and zero or one output symbol. This corresponds to the conventional definition of finite state transducers (FST) that are described in 2.3. The letter and the phoneme sequences are grouped into an equal number of segments. Such a grouping is called a joint segmentation. The joint probability is defined as:

$$p(\mathbf{g}, \varphi) = \sum_{\mathbf{q} \in S(\mathbf{g}, \varphi)} p(\mathbf{q}) \tag{2.8}$$

since there are many possible groupings of input sequences in general. The joint probability distribution $p(\mathbf{g}, \varphi)$ has thus been reduced to a probability distribution

$p(\mathbf{q})$ over graphone sequences $\mathbf{q} = q_1, \ldots, q_K$, which can be modeled using N-gram approximation:

$$p(q_1^K) \cong \prod_{j=1}^{K+1} p(q_j | q_{j-1}, \ldots, q_{j-M+1}) \qquad (2.9)$$

The probability distribution is then typically estimated by an Expectation-Maximization algorithm.

Generally, the problem of the wrong pronunciation in TTS is caused by a bad phonetic transcription. Traditional TTS systems are modular, one module's output is inputted into the next one. Because of this fact, the errors cumulate and thus the mistakes made by $g2p$ cannot be repaired. So if we want to improve the pronunciation, we can try to improve the $g2p$ as it is done in [**?**]. Authors in this work propose a method of exploiting a $g2p$ trained on a language with a high number of available resources to create a $g2p$ for language for which we do not have sufficient number of examples. This method relies on the existence of a conversion mapping between these two languages. Also it requires to do the conversion for every new language. In theory, a model can be created, that is able to transcribe grapheme sequences into an appropriate phonetic representation and can handle multiple languages. However, it needs to somehow obtain information, which language the input sequence comes from, which is not straightforwardly doable. This method has several drawbacks, because the language is not always known and the set of known languages is limited, so it does not realy solve the OOV problem. Also, it potentially requires a lot of training data. Moreover, if we want to learn a new pronunciation of just one word, it's more convenient to do it in a different way.

2.7.2 Learning pronunciation from spoken examples

. This group of methodologies aim to derive phonetic transcriptions directly from audio input. They are built on the theory of Automatic Speech Recognition which we discuss in 2.4. Authors of [**?**] introduce method of deriving correct pronunciation for a word in order to enlarge the recognizer's dictionary. They propose a data-driven approach to automatically add new words and their variants, respectively. The authors argue that in the spontaneous speech, the most frequent pronunciation does not need to be the one that is marked as correct and is used in the training phase. Thus the overall performance of the recognizer may be degraded since the phonetic units are bound with inadequate acoustics. The method is proposed, that relies on the use of both phoneme and word-level speech recognizer. The phoneme recognizer is constructed using smoothed bigram Language model. We discuss the phoneme recognizers in more detail in 2.4. The algorithm collects all occurences of words in the database and creates phonetic transcriptions using the recognizer. It then sorts the variants, rejects some of them an creates an $n - best$ list which is added to the dictionary. The recognizer can then be retrained, allowing multiple pronunciations for each word. Thus the recognition performance can be improved by an automatic procedure without the need for using the phonological rules.

In the work of McGraw et al. [6], the concept of Pronunciation Mixture Models

is introduced. The authors use a special kind of speech recognizer, the search space of which has four primary hierarchical components: a language model G, a phoneme lexicon L, phonological rules P that expand the phoneme pronunciations to their phone variations, and a mapping from phone sequences to context-dependent model labels C. All of the components metioned can be with advantage represented as FSTs and thus the full decoder network can be represented as a composition of these components: $R = C \circ P \circ L \circ G$. A probabilistic lexicon is considered, meaning, that several pronunciations are allowed for each word and there is no hard limitation that would force the recognzier to choose one. Instead, a kind of soft voting is considered, meaning, that each transcription is used with certain probability. Also, joint-sequence modelling is considered, as we introduced in 2.7.1, so each transcription is considered together with its orotgraphic form. That means, that we can describe the log-likehood of M utterances $D = \{\mathbf{u}_i, \mathbf{W}_i\}$ where $\mathbf{u}_i$ are speech data and $\mathbf{W}_i$ their transctiptions as follows:

$$\mathcal{L}(\Theta|D) = \sum_{i=1}^{M} log \sum_{B \in \mathcal{B}} P(\mathbf{u}_i, \mathbf{B}, \mathbf{W}_i; \Theta) \tag{2.10}$$

where $\mathbf{B}$ are respective phone sequences (i.e. pronunciations) and Θ represents the model parameters. Then, we derive using a chain rule:

$$P(\mathbf{u}_i, \mathbf{B}, \mathbf{W}_i; \Theta) = P(\mathbf{u}_i|\mathbf{B})P(\mathbf{B}|\mathbf{W}_i; \Theta)P(\mathbf{W}_i) \tag{2.11}$$

If we further assume, that pronunciation sub-units $\mathbf{b}_i \in \mathbf{B}$ are context independent, we can transcribe the above expression:

$$P(\mathbf{u}_i|\mathbf{B})P(\mathbf{B}|\mathbf{W}_i; \Theta)P(\mathbf{W}_i) = P(\mathbf{u}_i|\mathbf{B})(\prod_{j=1}^{k_i} P(\mathbf{b}_j|\mathbf{w}_j^i; \Theta))P(\mathbf{W}_i) \tag{2.12}$$

The model parameters are then estimated using the EM-algorithm. Parameters related to the language model, can be initialized with use of graphone language model. Several technical issues has to be dealt with, however the Pronunciation Mixture Models can be trained on the same data as traditional ASR engines and can be used to obtain phonetic transcriptions from audio data.

Similar approach is considered in the work of Reddy and Gouvea [10], except they do not have access to acoustic models or phone lattices. They use only the mistakes done by the recognizer as their source of information. An OOV word is passed through an ASR decoder giving an n-best word recognition output. Since the words are OOVs, every hypothesis will be a recognition mistake. These mistakes are then exploited, assuming that the following generative story of the recognition output for a word $\mathbf{w}$ holds:

1. A pronunciation baseform $\mathbf{b}$ is drawn from the distribution Θ.

2. A phonetic confusion function from the word $\mathbf{w}$ and the selected baseform $\mathbf{b}$ is applied in order to generate a phoneme sequence $\mathbf{p}$ with probability $P(\mathbf{p}|\mathbf{b}, \mathbf{w})$

3. A word sequence $\mathbf{e}$ with probability $P(\mathbf{e}|\mathbf{p}, \mathbf{b}, \mathbf{w}) = 1$ is generated using the pronunciation lexicon.

The authors model the joint probability of hypothesis and reference word $P(\mathbf{e}, \mathbf{w}) = P(\mathbf{w} \sum_b f_{e,b,w})$, where $f_{e,b,w} = P(\mathbf{e}|\mathbf{b}, \mathbf{w})$ is the phonetic confusion function and it is used to estimate the distribution $P(\mathbf{b}|\mathbf{w}, \mathbf{e})$. Thus they are able to derive pronunciations without access to ASR lattices, i.e. it only consideres the recognizer as a black-box.

2.7.3 Conclusion

Many approaches were introduced that are able to convert an utterance in ortographic or audio form to its phonetic representation. *G2P* converts the grapheme transcriptions, using only the text input. It is a well explored field of study with many different variants of realization. Although it achieves very good results nowadays, it suffers from the fact that same groups of letters may have different pronunciations in different languages. We typically don't have access to the information which language is considered and it may be difficult to get access to sufficient number of datasets. The latter problem can be partially solved by transfer learning as proposed in [**?**].

Alternatively, one can derive pronunciations directly from an audio signal. This approach has been also explored by some authors, however it usually requires quite low-level modifications of the speech recognizer. Also, the authors used the derived pronunciations to enlarge the phonetic dictionary of the recognizer, not to improve the Text-To-Speech sytems. We explore methods of merging the mentioned approaches to combine both textual and acoustic information and its usability when confronted with human judgments.

3. Proposed approaches

3.1 Basic overview

We aim to improve the pronunciation, that means, we need to derive phonetic transcriptions of a good quality. The straightforward way of achieving this is to use a grapheme-to-phoneme (g2p) module. As we discuss in 2.5 and 2.7, this approach has some limitations, mainly because the pronunciation of some tokens in different languages may vary. However, grapheme-to-phoneme conversion is a well explored area and its output may serve as a sufficiently good baseline. In section 2.7 we also discuss methods that are able to derive phonetic transcriptions directly from the speech signal, using a speech recognition framework. These methods are somewhat complex and they are typically specific for a particular recognizer. Rather than improving any of these methods, we aim to explore a possibility to combine them. That is, we want to exploit both textual and acoustic information to derive phonetic transcriptions of the desired words.

In the desired settings, the acoustic data with user's recordings are obtained online. We mean by this a situation. when the system asks the user for a correct pronunciation directly. It can be difficult to make user say words we are interested in, so the ultimate goal is to develop an algorithm able to extract those words from a dialogue history or ask user to say them in a not irritating way. The Crucial point of this approach is, that the user should not feel bored by the process. For the purpose of this work, we assume that we have already obtained recordings with correct pronunciations and work with it. In the real setting we would have to employ a dialogue policy to be able to gather our own recordings or process the history and try to identify respective words.

3.2 Used techniques

3.2.1 Phoneme recognition

To improve the pronunciation, it is essential to obtain information from the user's recording. We need to work on the phonetic level, so we have to gather phonetic transcription of the recording. In theory, we could use the standard output of the speech recognizer, i.e. a sequence of words and transcribe it phonetically. However, this is not desirable, because such procedure adds new source of potential errors, because we would have to use some mechanism, $g2p$ module for instance, that would create the transcription. Furthermore, since we aim to add new words to the TTS vocabulary, it is likely, that the word is missing also in the vocabulary of the speech recognition system. This fact means that the desired word will not appear in the recognizer's output and thus we cannot obtain any relevant transcription. Instead, we can modify the decoder used in the speech recognizer in such way, that it outputs sequences of phonemes rather than words. This can be done in several ways, some of which we introduce in the following text.

On the Automatic Speech Recognition (ASR) decoding process

We use the Kaldi framework for the ASR task and follow the procedure suggested in the official documentation. Kaldi constructs a Finite State Transducer for purposes of the decoding process[1]. This transducer is created as a composition of four other FSTs, $HCLG = H \circ C \circ L \circ G$. The meanings of the parts are as follows:

- G is an acceptor (i.e. its input and output symbols are the same) that encodes the grammar or language model.

- L is the lexicon; its output symbols are words and its input symbols are phones.

- C represents the context-dependency: its output symbols are phones and its input symbols represent context-dependent phones, i.e. windows of N phones

- H contains the acoustic model's definitions; its output symbols represent context-dependent phones and its input symbols reference the model.

As we can see, the G FST brings information contained in the language model to the decoding graph. We provide an example in Figure 3.2.1 which shows, that the L transducer determines the output elements. Hence, by modification of G and L we can change the output set of the decoded tokens.
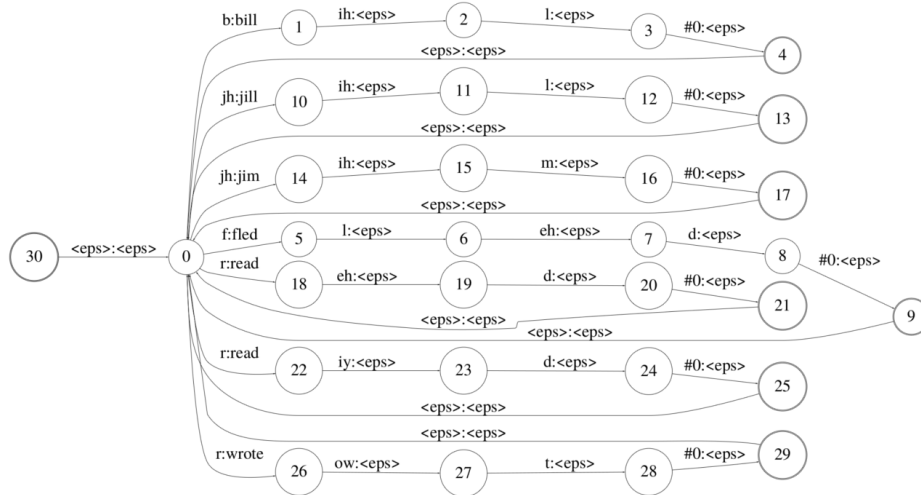
Figure 3.1: An example of the finite state transducer representing the lexicon[?]

Evaluation

In ASR systems, the most common phone recognition evaluation measures are Phone Error Rate (PER), or the related performance metric, phone accuracy rate. The latter is defined by the following expression:

$$Accracy = \frac{N_T - S - D - I}{N_T} \times 100\% \tag{3.1}$$

---

[1]http://kaldi-asr.org/doc/graph.html

where N_T is the total number of phonemes in the gold utterance and S, D and I are the substitution, deletion and insertion errors, respectively. From this perspective, we can see it as related to Levensthein distance of the gold and test sequences. Then, we can define $PER = 100\% - Accuracy$.

Another measure is correctness, which is similar to accuracy, but does not consider insertion errors. The number of insertion, deletion and substitution errors is computed using the best alignment between two token sequences: the manually aligned (gold) and the recognized (test). An alignment resulting from search strategies based on dynamic programming is normally used successfully for a large number of speech recognition tasks [**?**]. We use a *sclite* utility[2] to measure PER.

Modifying the Language Model

Language Model (LM) is an essential part of every ASR engine. It provides probabilities that a particular word occur in the considered context. For example we can imagine a model that assigns probability to every possible n-gram. The easiest approach is to use a so called 0-gram LM. It simply distributes the probability mass uniformly among the words.

Obviously, this model is not very good. If we want to employ prior knowledge about domain, we can use collected textual data and estimate probabilities of n-grams. If we want to decode phonemes, we can construct a language model that describes probabilities of phoneme occurrences instead of words.

We can then modify L and G transducers to change the output of the recognizer to phonemes. Ideally, we would use a set of name transcriptions to train the language model. However, it is difficult to get such corpus that would be also large enough to robustly estimate the LM parameters. So we create an artificial training dataset by transcribing written text with a *g2p* module. Alternatively, we could use phonetic alignments that are created during the decoding process and can be extracted from the recognizer. However, we did not choose this option, since the transcriptions that are obtained this way are not accurate enough, mainly because of the word recognition errors.

Creation of the decoding graph works with vocabulary and phonetic dictionary. The vocabulary contains a list of known words and the dictionary respective phonetic transcriptions. To change the transducers in the above mentioned way, we reduce the vocabulary to list of phonemes and setting the dictionary mapping to be identity. This way we force the decoder to output phonemes.

Timit dataset

The above approach has a disadvantage lying in the fact that it is difficult to obtain a language model of good quality without a proper corpus of transcriptions. This corpus may be found as a part of the timit dataset [**?**]. In fact, the Kaldi framework contains scripts, that train a model using this dataset. The timit dataset has got transcriptions labeled with granularity of individual phonemes. This means we can train the phonetic recognizer directly. Problem is that different set of phonemes is used but we can solve it by mapping the phonemes to IPA.

---

[2]http://www1.icsi.berkeley.edu/Speech/docs/sctk-1.2/sclite.htm

Worse is the fact, that the phone error rate measured on different data gives poor results. This can be caused by the fact that the dataset is not very big and thus the model can have problems accomodating to different acoustic conditions.

Modeling phoneme bigrams

The decoding graph is context dependent. It means that the decoder takes into account both the right and the left context of the phonemes. In order to model the pronunciation on various positions in the word, each phoneme is represented in several variants, depending in which part of the word it is located ($beginning, end, inside$). One option is to model each phoneme as a singleton as we have done in previous. Alternatively, we can create a decoding graph, that models bigrams of phonemes. This way the information about the context is preserved. Moreover, we can model the input more accurately and distinguish pauses between words from space between phonemes. This approach proved to yield the best results.

Exploring the speech recognizer

The decoding process in Automatic Speech Recognition yields a list of hypotheses, ordered by their confidence scores, so called n-best list. Usually, the best hypothesis is chosen and used since it is the most relevant result. However, interesting information may be found deeper of the list. We explore this theory in the following paragraphs.

In section 3.2.1 we have introduced Phone Error Rate (PER). We extend the term here and define **Oracle Phone Error Rate**:

$$OPER = \min_{h \in hypotheses} PER(h, gold) \tag{3.2}$$

We can see that Oracle Phone Error Rate is PER of the best hypothesis (if we consider the hypothesis with the lowest PER as the best one.) In real setup we obviously does not have such oracle that would tell us the best hypothesis, so we have no other option but looking at the top of the n-best list, computing PER per each hypothesis and choosing the closest one. Hence we approximate the true Oracle PER, since if we explore "sufficiently many" best alternatives, the probability that we did not see the best one is very low. We give an overview of our results in Figure 3.2.1.

Using Oracle Phone Error Rate to evaluate the recognizer is basically cheating, since in real decoding process, we are not given the gold transcription which we could use to pick the best hypothesis. Nevertheless, it indicates some properties of the decoder, namely that the hypothesis located at the top is not always the best one. We explore this phenomenon further by looking at the particular position of the closest (best) hypothesis. In Figure 3.2.1 we plot the histogram showing, number of times the closest hypothesis appeared in particular depth. We can see, that relevant hypotheses can be found in bigger depth.
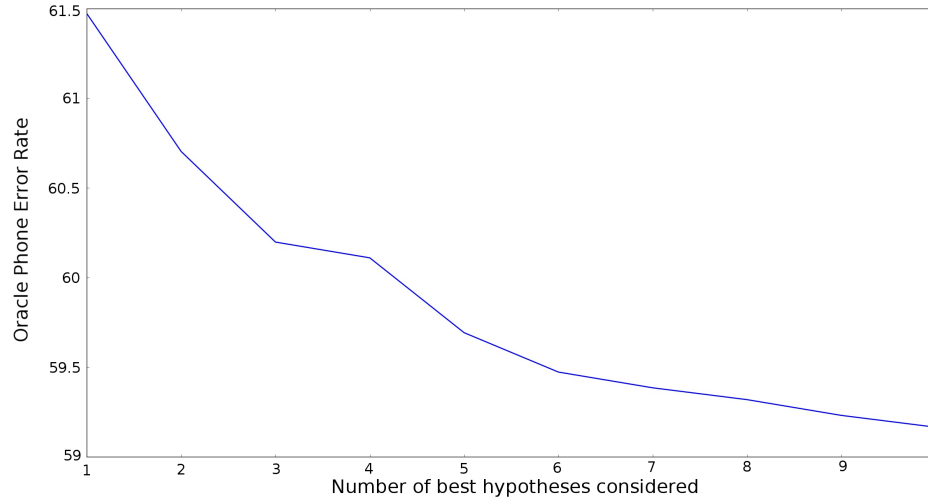
Figure 3.2: A plot of Oracle Phone Error Rate dependency on the depth of the n-best list we explore. We can see, that the Oracle PER decreases with the depth, so it makes sense to try to exploit the hypotheses located deeper.
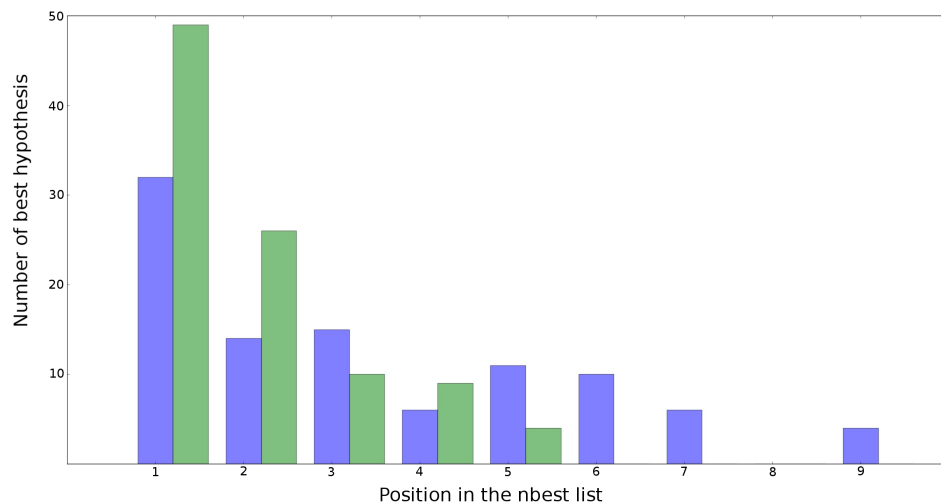

Figure 3.3: Histogram that visualizes the distribution of the hypothesis, that is the closest to the gold reference utterance in terms of Phone Error Rate. The bars show the number of times the best hypothesis was found on the particular position. Two settings are included in this figure: the green bars describe situation, when first five positions were used, the blue ones were measured with nine positions

3.3 Text-to-speech systems' evaluation

The TTS evaluation[**?**] is a desired task since we often want to choose between two or more engines. However the task introduces several challenges some of which we mention in this section. We need to define a metric in order to evaluate the TTS. Such a metric will be dependend on particular task's setting and may consist of several variables. We can describe a taxonomy of different evaluation methods, depending from which point of view we look at it we distinguish

between:

- *Black box vs. Glass box evaluation* - whether we evaluate the TTS system as a whole or look at particular modules of it.

- *Human vs. automated* There are two fundamentally distinct ways of evaluating speech synthesizers: one is to use human subjects; the other to automate the evaluation process. Nowadays, it is necessary to employ humans in the evaluation process, mainly to evaluate the quality of integration and functionality of the system.

- *Global vs. analytic* approach tell us whether we rate global aspects such as naturalness and overall quality or more analytic aspects like vowel and consonant clarity, tempo or appropriateness of stresses and accents.

A critical measurement of a TTS system is whether or not human listeners can understand the text read by the system. So called *intelligibility* tests were developed to measure this aspect. For example Diagnostic Rhyme Test is used that uses set composed of pairs of similair words. Those are synthesized with a TTS engine and the listeners try to correctly distinguish between them. The intelligibility is the evaluated and compared to listening human speakers. An example set of words that may be used for such test is given in Figure 3.3.

Voicing		Nasality		Sustenation		Sibilation		Graveness		Compactness	
veal	feel	meat	beat	vee	bee	zee	thee	weed	reed	yield	wield
bean	peen	need	deed	sheet	cheat	cheep	keep	peak	teak	key	tea
gin	chin	mitt	bit	vill	bill	jilt	gilt	bid	did	hit	fit
dint	tint	nip	dip	thick	tick	sing	thing	fin	thin	gill	dill
zoo	sue	moot	boot	foo	pooh	juice	goose	moon	noon	coop	poop
dune	tune	news	dues	shoes	choose	chew	coo	pool	tool	you	rue
vole	foal	moan	bone	those	doze	joe	go	bowl	dole	ghost	boast
goat	coat	note	dote	though	dough	sole	thole	fore	thor	show	so
zed	said	mend	bend	then	den	jest	guest	met	net	keg	peg
dense	tense	neck	deck	fence	pence	chair	care	pent	tent	yen	wren
vast	fast	mad	bad	than	dan	jab	gab	bank	dank	gat	bat
gaff	calf	nab	dab	shad	chad	sank	thank	fad	thad	shag	sag
vault	fault	moss	boss	thong	tong	jaws	gauze	fought	thought	yawl	wall
daunt	taunt	gnaw	daw	shaw	chaw	saw	thaw	bong	dong	caught	thought
jock	chock	mom	bomb	von	bon	jot	got	wad	rod	hop	fop
bond	pond	knock	dock	vox	box	chop	cop	pot	tot	got	dot

Figure 3.4: An example set of words used to perform the Diagnostic Rhyme Test

While a TTS system has to be intelligible, this does not guarantee user acceptance, because its quality may be far from that of a human speaker. That is the reason, why methods like Mean Opinion Score (MOS) are frequently used. MOS is human-subject judgment testing that asks the subjects to rate several samples on the scale of 1 to 5 (1 = Bad, 2 = Poor, 3 = Fair, 4 = Good, 5 = Excellent). The scores are averaged, resulting in an overall MOS rating. We have used this method of labeling to annotate our sample dataset.

Normalized MOS scores for different TTS systems can be obtained without any preference judgments. Nevertheless, sometimes comparisons are desired, especially for systems that are informally judged to be fairly close in quality. So

called Category Rating (CCR) method, may be used. In this method, listeners are presented with a pair of speech samples on each trial. The order of the system A system B samples is chosen at random for each trial, providing, that half of the trials, the system A sample is followed by the system B and the order is reversed in the remaining trials. Listeners rate the pair with labels meaning, that the second utterance is:

- 3 - much better

- 2 - better

- 1 slightly better

- 0 about the same

- -1 slightly worse

- -2 worse

- -3 much worse

Sometimes the granularity can be reduced as much as simply "prefer A/prefer B". We use this method to tell which transcription is better.

4. Experiments

4.1 Data

Although quite a lot of speech datasets have been published, we were interested in datasets that would contain a lot of OOV words. We mean by this, that it contains a lot of words, that are not present in a vocabulary of the speech recognition and text-to-speech that we have used. We aim for this, because we would like to explore such words. They should be difficult to pronounce for the text-to-speech engine and thus there is space for an improvement.

4.1.1 Dataset D_1 - artificial

Dataset D1 contains 100 randomly chosen Dutch Names. Their recordings were used to derive phonetic transcriptions and we created three other transcriptions, as we describe further in section TODO. Annotators first listened each word read by native speaker and then listened the synthesized recordings and rated it with a label one to five.

4.1.2 Dataset D_2 - artificial

D2 is also manually annotated set of 110 words. 70 words are English, remaining 40 are Czech. Czech words were chosen to contain lot of diacritic hence to be difficult to pronounce for TTS trained on English. Each word was synthesized by four engines described in the beginning of this section. Each recording was then manually annotated. **Data labeling**. Labeling of 2_1 was performed by three annotators. Each of them had to listen each of the words three times, each time synthesized by different engine. He then labeled each of the recording with a discrete number ranging from one to five. The recordings labels were averaged. Thus we obtain labels per each recording and we have three labels in total for each word. We average these three labels and obtain a score between one and five, that describes overall quality of the word's pronunciation.

4.1.3 Dataset D_3 - Autonomata Spoken Name Corpus

The Autonomata Spoken Names Corpus [**?**] is a database of approximately 5.000 read-aloud first names, surnames, street names, city names and control words. It was published by Dutch-Flemish HLT Agency[1]. The corpus consists of a Dutch part and a Flemish part. Besides the speech data, the corpus also comprises phonetic transcriptions, speaker information and lists containing the orthography of the read names. Spoken utterances of 240 speakers living in the Netherlands (NL) or in Flanders (FL) are included in the corpus. Since we have both ortographic and phonetic transcriptions available in addition to the audio recording, the dataset is ideal for our purposes, because we can evaluate the quality of our transcriptions.

---

[1]http://tst-centrale.org/nl/

Annotators	Rec. 1	Rec. 2	Rec. 3	Rec. 4	mean
A_1 **vs** A_2	0.51	0.30	0.47	0.63	**0.48**
A_1 **vs** A_3	0.44	0.44	0.34	0.59	**0.45**
A_2 **vs** A_3	0.48	0.34	0.38	0.57	**0.44**
mean	**0.48**	**0.36**	**0.40**	**0.60**	

Table 4.1: The table contains results of linear regression model training on aggregated MFCC vectors. Correlation with human judgment and R^2 measure for respective cross-validation folds are shown.

4.1.4 Annotation procedure

We have created some datasets (D_1, D_2) ourselves, so we had to obtain annotations for it. Hence, we asked three annotators to label the recordings. They were given instructions to judge the recordings with respect to quality of the pronunciation, not the sound quality of the recording . However, the criterions were defined quite vaguely, that is they were not told which aspects of pronunciation they should focus on. Therefore, the ratings were quite diverse and they were subjective according to respective annotators' preferences.

Because of that, we have explored how individual annotators' ratings correlate. We sum up the results in 4.1.4

4.2 Identifying difficult words

If we want to improve pronunciation of the synthesized recordings, we should first identify words that are mispronounced. Obviously, we could let the TTS system pronounce each word in a best effort style and the user would identify mistakes and correct them. This approach hase several drawbacks. First, the communication with the user is rather complicated, since the incorrectly pronounced word has to be isolated prior to obtaining the correct pronunciation. Second, it may be unpleasant for the user if he or she has to undergo this process and hear the incorrect pronunciation. Thus, it would be better if we could recognize the possibly difficult words somehow. It would mean that we can ask the user directly for the correct pronunciation of the word. We explore this issue in the rest of this section.

4.2.1 Measuring the difficulty

In order to estimate the difficulty of each word's pronunciation, we propose three measures in this section. The key idea is that we obtain values for each of this measures and then combine them in one feature vector that represents the recording. Then we can train a classifier that learns to predict quality of the pronunciation.

M_1 measure - averaged Mel Cepstral Distortion

We discuss Mel Cepstral Distortion in detail in 2.2. We use it here to define the M_1 measure. We can describe it as follows:

$$M_1(k) = \frac{1}{N} \sum_{(i,j)} MCD(r_{ki}, r_{kj}) \tag{4.1}$$

Where $N = \binom{3}{2}$, (i,j) stands for every combination of synthesizers, r_{ki} is name k synthesized by engine i and MCD is the Mel Cepstral Distortion. The key idea motivating this measure is that if there is a problem with a pronunciation of some part of the word, every synthesizer has to deal with it somehow. It is likely, that each of them will do it in sligthly different way. This implies that the pairwise MCD will increase.

Nevertheless, the process of computing the value of M_1 introduces some problems. The troubles stem from the fact, that there is a high variability in the data obtained from the synthesizers. That is, one synthesizer's output may differ greatly from the others, so it biases the result. We can see the effect in Figure 4.2.1. If we had enough synthetizers, we could afford to ignore values of the outliers and thus smooth the results.

Figure 4.1: An example demonstrating, how the measured M_1 between two recordings can vary. Red and blue bars represent respective recordings, their height corresponds to the measured M_1. On the left, there are measurings using only two synthesizers ($\sim MCD$), in the betwwen, there are triples and the most right column corresponds to four synthesizers. One outlying example can negatively influence the relevancy of the measure. While the recording represented by the red columns has quite similar MCD between every pair, the other one has outliers and thus the resulting value has got low confidence.

We tested, how good is the correlation of M_1 measure and human judgement. First, we synthesized the words by different synthesizers. Each recording created this way was then labeled by a human. Set of recordings for each word was

then used to compute M_1 value and this was compared with the mean of the respective human judgements. Results on our sample dataset with the use of the 0-th coefficient and without it are shown in Figure 4.2.1. Based on these results, it may seem that the M_1 measure is not very good, the best correlation was 0.471. However, it may add a useful piece of information to the feature vector.

Figure 4.2: Plot of the M_1 measure correlation with human judgments. Points are horizontally spaced and colored according to gold labels. The vertical axis shows M_1 value. The blue lines represents estimate yielded by least square fit. The top figure includes the 0-th mel cepstral coefficient, while the figure on the bottom does not.

The MCD has one property that can be looked at as a disadvantage: It does not weight its coefficients, more accurately it weights all with the same weights. We propose to change this - for example a linear model could be trained, that finds the combination of MFCC's that corresponds the best. The coefficients derived by this model would be difficult to interpret without further research. However, we can train it and try to use it instead of typical MCD. The disadvantage of this approach is that we need labeled data, while the previous method works in an unsupervised way.

We want to supply the computation of M_1 by a linear model prediction. To

Fold	Correlation	R^2
1	0.77	0.74
2	0.03	0.75
3	0.82	0.76
4	0.35	0.77
5	0.69	0.74
6	0.53	0.72
7	0.94	0.75
8	0.68	0.73
Mean	**0.60**	**0.75**

Table 4.2: The table contains results of linear regression model training on aggregated MFCC vectors. Correlation with human judgment and R^2 measure for respective cross-validation folds are shown.

train such a model, we must first prepare the data. We use the D_1 dataset described in section 4.1. Since we want to use a simple linear model, we need to work with fixed size vectors. Nevertheless, the input data are variable-length sequences of vectors. We first transform each synthesized recording by summing all its mel cepstral coefficients. Thus we get N M-sized vectors, where N is the number of synthesizers used and M is the order of mel cepstral analysis. We then compute differences in each position between every pair of recognizers and sum those to obtain one resulting vector of length M. We use labels obtained from annotators normalized to interval $[0, 1]$ as a target variable. We then evaluate the model using 8-fold cross-validation. The results are shown in Figure 4.2.1. The correlations and R^2 for respective folds are shown in Table 4.2.1.

We can train a model that works with untransformed data representation, that is all the full sequences. The input has a big dimension, however, the number of parametres is not that big, because we want to train only one scale factor per one position in the MFCC vector. TODO

M_2 measure - averaged phonetic distance

Another measure we can possibly use is based on phonetic transcriptions. We first recognize the recordings with a phoneme recognizer, thus we obtain a sequence of characters per each recording. We can compute pairwise distances of these transcriptions, using for example Levensthein [?] or Hamming distance, which we normalize. The motivation is similar to the M_1 measure. Assuming the difficult words have positions that are problematic for the TTS engine, the recognized phonemes on these positions should differ and thus the distance between transcriptions should increase. Four our purposes, we have used Levenshtein distance as a metric.

The M2 measure is described by equation:

$$M_2(k) = \frac{1}{N} \sum_{(i,j)} \frac{LD(Hyp(r_{ki}), Hyp(r_{kj}))}{max(len(r_{ki}, len(r_{kj})))} \tag{4.2}$$

Where $N = \binom{3}{2}$, (i, j) stands for every combination of synthesizers and r_{ki} is name
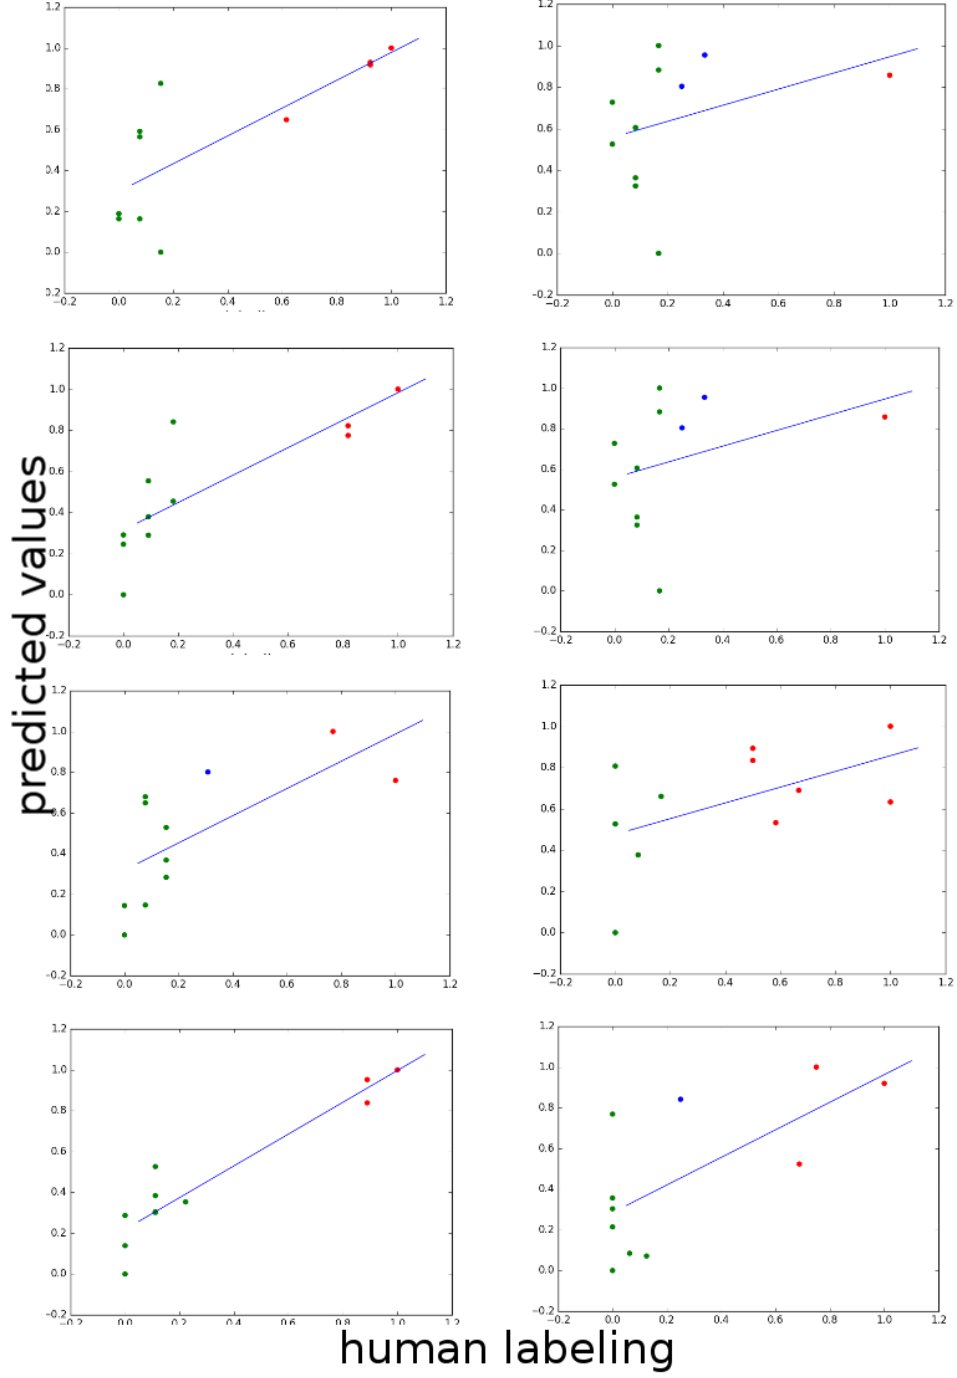
Figure 4.3: Plots of the correlation of the measure computed by model in cross-validation folds. The predicted values are on the vertical axis, the horizontal position represents the human judgments. The blue lines shows least-squares fit. Although the training data are rather small, it shows, that the model can output meaningful values.

$\{k\}$ synthesized by engine $\{i\}$. *LD* means Levenshtein Distance and *Hyp* represents the best hypothesis from the phonetic recognizer. Note, that we normalize the distance by length of the longer transcription.

Figure 4.4: Scatterplot showing relation between the M_1 and M_2 measures. We can see, that the measures are rather uncorrelated. However, they can still be complementary.

M_3 measure - occurence of bad bigrams

The M_3 measure is based on a different approach. We want to learn the typical mistakes of a grapheme-to-phoneme converter, more specificaly we want to explore which groups of graphemes are difficult to transcribe for the $g2p$. We suppose that words with many occurences of such groups are difficult to pronounce. To compute the measure, we need two corpuses (C_1 and C_2) in different languages and a $grapheme-to-phoneme$ ($g2p$) training algorithm. The corpuses should consist of pairs (word, transcription). We can then prepare for computation of M_3 in three stages:

1. Train $g2p$ model G_1 on corpus C_1 in language L_1.

2. Use trained G_1 to transcribe words contained in corpus C_2.

3. Identify problematic parts.

Stage 3 needs further description. We obtained list of triples, i^{th} of which is structured:

(word w_i from C_2, original transcription t_i^o, transcription t_i^h derived by G_1)

We can then use the Dynamic Time Warping algorithm to obtain pairwise alignments of these sequences, illustrated on Figure 4.2.1 In fact, we first transform the phoneme sequences into sequences of bigrams. This is because graphemes and phonemes are not in one-to-one correspondence and bigrams capture the relations better. Once we have the alignments, we can identify positions, in which the original transcriptions from C_2 differ from the hypothesis derived in Stage 2. Respective bigrams from the original word can then be identified. For each bigram, we count number of times it was marked as difficult. Each word can then be scored according to number of bad bigrams it contains. To evaluate the M_3 measure, the MaryTTS framework was used, because we can extract phonetic transcriptions from it and thus use its $g2p$ module. This means that we

Figure 4.5: Sample alignment of two phonetic sequences

are able to derive the transcriptions with exactly the same module that is used in the synthesizer. We have also used the Cereproc engine. Because we do not have access to the *g2p* that Cereproc uses, we trained our own model on the CMU dictionary[2]. The data was then processed and the confusion matrix was created, containing number of times, respective bigrams are confused with each other. The matrix turns out to be quite sparse, which is not surprising, since the majority of bigram pairs are interchanged with probability nearly zero.

If we restrict the matrix only to values greater than a certain threshold, its dimensions decrease dramatically and we can visualize it; the result can be seen in Figure 4.2.1 The scoring procedure counts for each word a number of occurrences of bigrams that were confused and the number of confusions is summed and divided by the length of the word. We plot the scores in Figure 4.2.1. The shape of the curve corresponds to the fact that many bigrams do not occur a lot, or are not problematic.

The correlation of this measure and annotations was 0.20 for the MaryTTS and 0.31 for the Cereproc, respectively. Intuitively, if we have access to the *g2p* module, which is the MaryTTS case, it is expectable, that the results would be better. However, this turns out not to be true, which is surprising. TODO: why?

The disadvantage of this approach is that it relies quite heavily on the alignment process, which is imperfect, so sometimes it can mark as confused pair of bigrams, that is not correct. Also id occurences of some bigrams are very sparse, then its scores may be estimated poorly. Another possible problem is, that it relies on the corpuses with phonetic transcriptions and then it is specific to certain pair of languages.

Measure combination

We have seen that neither of the measures correlates well with the annotators' labeling. Despite this fact, it may by interesting to explore, whether the combination of the measures can work better. As described in the introduction to this section, we combine the measured values to one feature vector that represents the word. We then train a *Ridge Regression* model using the averaged human labels as a target variable. We have used 5 fold cross validation and measured the R^2 and $MeanSquaredError(MSE)$, which we averaged over all the folds. The measured MSE was 0.43 and R^2 was 0.58. We remind, that the human labels are discrete values in the interval $[1, 5]$, so the MSE value is good, saying that we differ from the human label by less than 0.5.

---

[2]http://www.speech.cs.cmu.edu/cgi-bin/cmudict

Figure 4.6: Frequencies of confusing respective bigrams, i.e. how many times it occurs in some confused pair. We can see, that the majority of bigrams has very low frequencies.

We can choose a threshold and state, that if the model's predicted value is greater than the threshold, the respective word is hard to pronounce and we should try to improve the pronunciation. If we plot the results, as in Figure 4.2.1, we can see, that setting the threshold to value 2.5, which is exactly in the middle of the scale is reasonable choice. Reformulating the problem as a binary classification task gives us another insight. We thus have to decide if a given example is pronounced correctly or not. We aim to identify incorrectly pronounced words and we refer to these words as positive examples. Thus we can compute the *precision and recall* [3] and obtain the $F_1 - score$[4] 0.93. Based on these experiments, our method has reasonable behavior and may be used to identify difficult recordings confidently.

Discussion

We proposed and explored three measures for a purpose of recognition of words that are difficult to pronounce for a TTS system. Our goal is to develop a method that would identify such words before they are introduced to the user. Exploration of these measures showed, that they does not correlate well with the human judgments. Nevertheless, the experiments indicate, that the measures are complementary in the sense, that if we combine them in a feature vector, we can train a model that predicts reasonable values. If we determine a threshold, we can use this model for classification purposes. Although the results are promising,

---

[3]https://en.wikipedia.org/wiki/Precision_and_recall
[4]https://en.wikipedia.org/wiki/F1_score

Figure 4.7: Visualization of the confusion matrix computed from the results.

the computation process that preceeds the creation of the feature vector is not very convenient. We have to have access to speech engines, non-trivial phonetic corpuses (M_3) and phoneme recognizer (M_2).

Figure 4.8: Plot of the labels given by humans (vertical axis) and predicted values (horizontal axis). The vertical line represents the threshold with respect to the predictions, the horizontal dashed line ilustrates true division.

Conclusion

Bibliography

[1] Arpabet overview. https://nlp.stanford.edu/courses/lsa352/arpabet.html. Accessed: 2017-04-16.

[2] Maximilian Bisani and Hermann Ney. Joint-sequence models for grapheme-to-phoneme conversion. *Speech Communication*, 50(5):434–451, 2008.

[3] Michael J Dedina and Howard C Nusbaum. Pronounce: a program for pronunciation by analogy. *Computer speech & language*, 5(1):55–64, 1991.

[4] Ronald M Kaplan and Martin Kay. Regular models of phonological rule systems. *Computational linguistics*, 20(3):331–378, 1994.

[5] John Lucassen and Robert Mercer. An information theoretic approach to the automatic determination of phonemic baseforms. In *Acoustics, Speech, and Signal Processing, IEEE International Conference on ICASSP'84.*, volume 9, pages 304–307. IEEE, 1984.

[6] Ian McGraw, Ibrahim Badr, and James R Glass. Learning lexicons from speech using a pronunciation mixture model. *IEEE Transactions on Audio, Speech, and Language Processing*, 21(2):357–366, 2013.

[7] Mehryar Mohri. Finite-state transducers in language and speech processing. *Computational linguistics*, 23(2):269–311, 1997.

[8] Mehryar Mohri. Weighted automata algorithms. In *Handbook of weighted automata*, pages 213–254. Springer, 2009.

[9] Chotirat Ann Ratanamahatana and Eamonn Keogh. Everything you know about dynamic time warping is wrong. In *Third Workshop on Mining Temporal and Sequential Data*. Citeseer, 2004.

[10] Sravana Reddy and Evandro B Gouvêa. Learning from mistakes: Expanding pronunciation lexicons using word recognition errors. In *INTERSPEECH*, pages 533–536, 2011.

[11] Tim Schlippe, Sebastian Ochs, and Tanja Schultz. Grapheme-to-phoneme model generation for indo-european languages. In *Acoustics, Speech and Signal Processing (ICASSP), 2012 IEEE International Conference on*, pages 4801–4804. IEEE, 2012.

[12] Terrence J Sejnowski and Charles R Rosenberg. *NETtalk: A parallel network that learns to read aloud*. MIT Press, 1988.

[13] P. Taylor. *Text-to-Speech Synthesis*. Cambridge University Press, 2009. ISBN 9781139477260. URL https://books.google.cz/books?id=TOO-NHZx7kIC.

[14] Paul Taylor and Amy Isard. Ssml: A speech synthesis markup language. *Speech communication*, 21(1-2):123–133, 1997.

[15] Aäron van den Oord, Sander Dieleman, Heiga Zen, Karen Simonyan, Oriol Vinyals, Alex Graves, Nal Kalchbrenner, Andrew Senior, and Koray Kavukcuoglu. Wavenet: A generative model for raw audio. *CoRR abs/1609.03499*, 2016.

[16] Yuxuan Wang, RJ Skerry-Ryan, Daisy Stanton, Yonghui Wu, Ron J Weiss, Navdeep Jaitly, Zongheng Yang, Ying Xiao, Zhifeng Chen, Samy Bengio, et al. Tacotron: A fully end-to-end text-to-speech synthesis model. *arXiv preprint arXiv:1703.10135*, 2017.

[17] Kaisheng Yao and Geoffrey Zweig. Sequence-to-sequence neural net models for grapheme-to-phoneme conversion. *arXiv preprint arXiv:1506.00196*, 2015.

List of Figures

List of Tables

List of Abbreviations

Attachments